

Refund Policy Q&A Application Documentation

1. Overview

The **Refund Policy Q&A** application is a Streamlit-based interactive tool that allows users to upload a PDF document containing their company refund policies and then ask questions about its content. The system uses Natural Language Processing (NLP) techniques leveraging LangChain, HuggingFace embeddings, and a transformer-based language model (Google Flan-T5) to provide relevant answers extracted from the uploaded document.

2. Features

- Upload company policy PDF files
 - Extract text from PDF pages automatically
 - Split text into chunks for better semantic search
 - Create or load a cached FAISS vector store to speed up embedding search
 - Use sentence-transformer embeddings for document indexing
 - Query the document with natural language questions
 - Get relevant answers generated by a language model
-

3. Installation & Setup

Requirements

- Python 3.7+
- Streamlit
- PyPDF2
- LangChain
- FAISS
- HuggingFace transformers and sentence-transformers
- Additional dependencies (listed in requirements.txt)

Install dependencies

```
bash

pip install streamlit PyPDF2 langchain faiss-cpu transformers sentence-transformers
```

Running the app

```
bash

streamlit run app.py
```

Open your browser and navigate to the URL shown (usually <http://localhost:8501>).

4. Code Explanation

Key components:

- **PDF Upload & Extraction**

Uses `st.file_uploader` to accept PDF files, and `PyPDF2` to extract text from each page.

- **Text Splitting**

Uses `LangChain's CharacterTextSplitter` to break large text into manageable chunks (1000 chars with 200 overlap).

- **Embeddings**

Utilizes `HuggingFace's "sentence-transformers/all-MiniLM-L6-v2"` model to embed the text chunks.

- **Vector Store**

Uses `FAISS` to build or load a local vector index for similarity search, keyed by an MD5 hash of the uploaded PDF file contents.

- **Retriever & QA Chain**

The retriever fetches top-k relevant chunks, then passes them along with the user question to the LLM (`Google Flan-T5` via `HuggingFace` pipeline) wrapped in `LangChain's RetrievalQA`.

- **User Interface**

Streamlit displays the upload widget, input box for questions, and the model's answers.

5. How to Use

1. Upload a PDF file containing your company's refund policy.
 2. Wait for the text extraction and vector store indexing (cached for repeated uploads).
 3. Enter a question about the refund policy in the text input box.
 4. Receive an answer based on the uploaded document.
 5. Repeat with new questions or upload another document.
-

6. Sample Inputs and Outputs

Example 1

Uploaded file: Refund_Policy_of_ipu.pdf (0.8MB)

User question:

What happens if a candidate applies for withdrawal after the last date of admission in medical courses?

Answer:

he/she can apply for withdrawal within the specified date, an amount of Rs. 1000/- shall be deducted and the balance amount shall be refunded

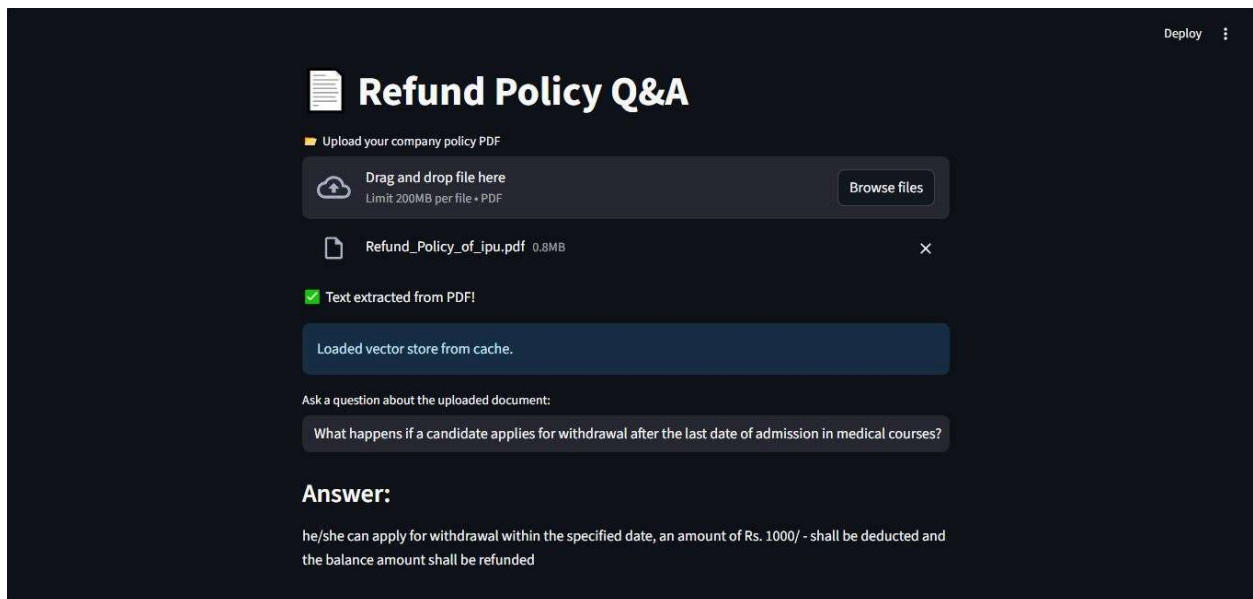


Fig1: Example1 Snapshot

Example 2

Uploaded file: Refund_Policy_of_ifinltd.pdf (61.4KB)

User question:

What are additional non-returnable items?

Answer:

Gift cards • Downloadable software products • Some health and personal care items

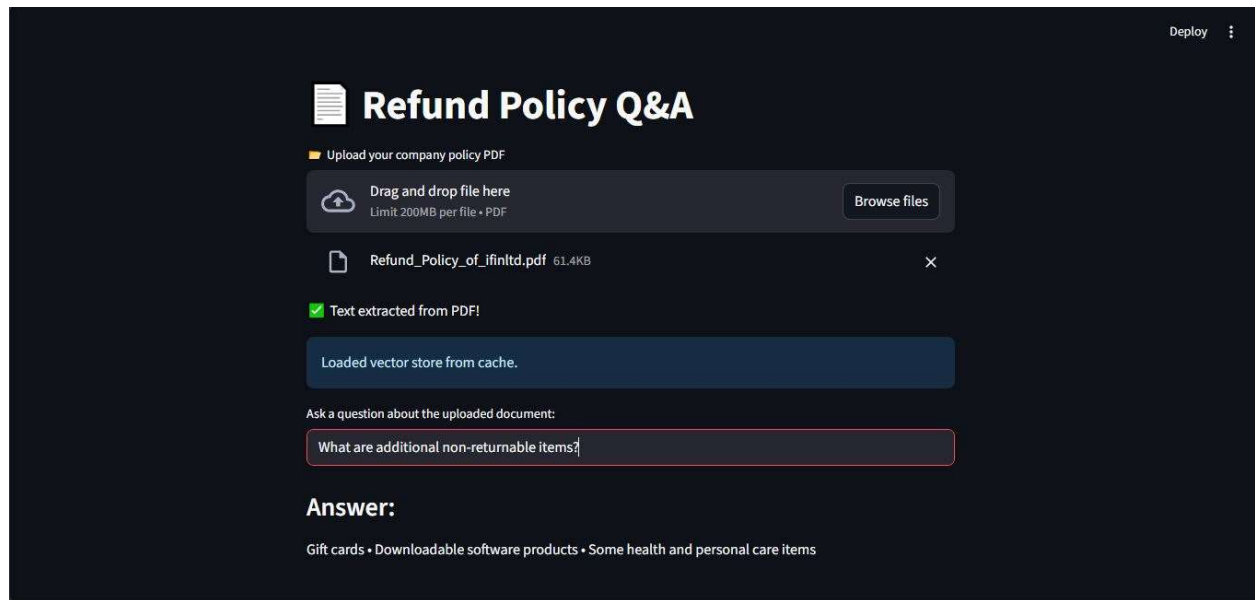


Fig2: Example2 Snapshot

7. Notes & Improvements

- **Caching:** The vector store is cached based on an MD5 hash of the uploaded file bytes, improving speed on repeated uploads of the same document.
- **Answer Quality:** Answers depend on the quality and scope of the uploaded PDF and the LLM used.
- **Extensions:**
 - Add multi-document upload and search.
 - Improve chunking strategy.
 - Experiment with larger or more powerful LLMs.
 - Highlight source text snippets in the answer.

- Add logging and error handling for robustness.

8. Code Listing

```
RetrievalQA_A4.py X
RetrievalQA_A4.py > ...
2 import streamlit as st
3 from langchain.chains import RetrievalQA
4 from langchain.vectorstores import FAISS
5 from langchain.embeddings import HuggingFaceEmbeddings
6 from langchain.llms import HuggingFacePipeline
7 from langchain.text_splitter import CharacterTextSplitter
8 import PyPDF2
9 import os
10 import hashlib
11 from transformers import pipeline
12
13 st.title("Refund Policy Q&A")
14
15 uploaded_file = st.file_uploader("Upload your company policy PDF", type="pdf")
16
17 if uploaded_file is not None:
18     pdf_reader = PyPDF2.PdfReader(uploaded_file)
19     full_text = ""
20     for page in pdf_reader.pages:
21         text = page.extract_text()
22         if text:
23             full_text += text + "\n"
24
25     st.write("Text extracted from PDF!")
26
27     file_bytes = uploaded_file.getvalue()
28     doc_hash = hashlib.md5(file_bytes).hexdigest()
29     index_path = f"faiss_index_{doc_hash}"
30
31     text_splitter = CharacterTextSplitter(separator="\n", chunk_size=1000, chunk_overlap=200)
32     texts = text_splitter.split_text(full_text)
33
```

Fig3: Input Code Part1 Snapshot

```
RetrievalQA_A4.py X
RetrievalQA_A4.py > ...
33
34 embeddings = HuggingFaceEmbeddings(model_name="sentence-transformers/all-MiniLM-L6-v2")
35
36 if os.path.exists(index_path):
37     # Allow loading pickle safely because it's your own file
38     vector_store = FAISS.load_local(index_path, embeddings, allow_dangerous_deserialization=True)
39     st.info("Loaded vector store from cache.")
40 else:
41     vector_store = FAISS.from_texts(texts, embeddings)
42     vector_store.save_local(index_path)
43     st.info("Loaded cached vector store.\n Created new vector store and cached it.")
44
45 retriever = vector_store.as_retriever(search_type="similarity", search_kwargs={"k":3})
46
47 pipe = pipeline("text2text-generation", model="google/flan-t5-small")
48 llm = HuggingFacePipeline(pipeline=pipe)
49
50 qa = RetrievalQA.from_chain_type(llm=llm, retriever=retriever)
51
52 question = st.text_input("Ask a question about the uploaded document:")
53
54 if question:
55     with st.spinner("Finding answer..."):
56         answer = qa.run(question)
57         st.markdown("### Answer:")
58         st.write(answer)
59 else:
60     st.info("Please upload a PDF file to get started.")
61
```

Fig4: Input Code Part2 Snapshot

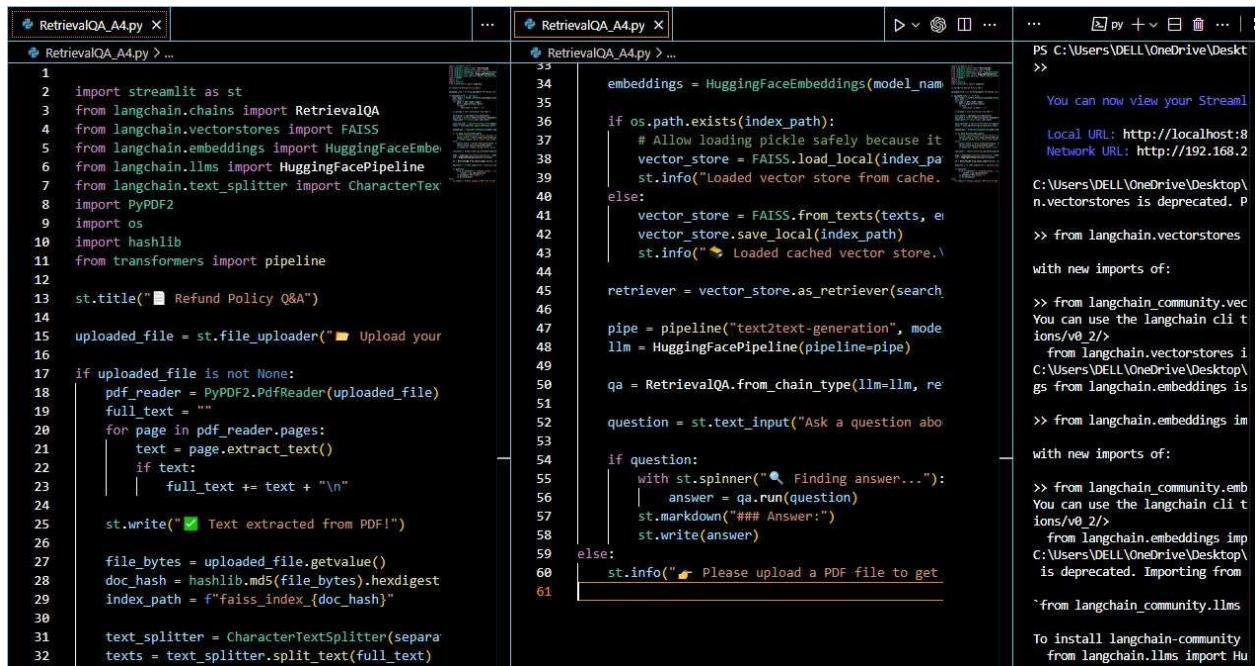


Fig5: Input Code Snapshot